
Evaluation of Graph Neural Network on Simple Weather Forecast Task

Ambrose Lee¹

Abstract

Weather forecast has been playing a critical role in contemporary society, and improving the accuracy of temperature forecast could bring numerous social and economic benefits. In this work, our objective is to evaluate different machine learning techniques for temperature prediction, and these results are going to serve as the foundation for upcoming more complicated weather forecast task. We experimented with different machine learning models and graph neural networks to predict the mean temperature one day after. We evaluated the accuracy of different models, and how number of layers and batch size affect the performance of the model. Code is available at <https://github.com/ambr0se/TemperatureForecast>.

1. Introduction

Weather forecasts have played an important role in modern society, not only for normal people but also to various industries. Governments rely on long-term temperature forecast to devise climate policy, while the agricultural sector plan and revise their farming strategies according to short-term temperature forecast (Lobell & Field, 2007). A more accurate temperature could bring immense economic and social benefits, incentivizing researchers to investigate the field.

Recently, graph neural networks (GNNs) have gained traction in machine learning community, and have been applied to various tasks, most commonly node classification and link predictions (Zhou et al., 2018). Designed to process graph-structured data, GNN excels in a wide range of application, but some studies found that GNNs are subject to limited representation power, and many model architectures and kernels fail to surpass 1-WL test (Xu et al., 2018).

This paper serves as a foundation work of predicting weather

with graph neural networks. Specifically, we examined models and techniques that perform well on the dataset, and these results would serve as the corner stone of the coming more specific and complicated weather forecast task. Through this experiment, we aimed to study that 1) despite the limited representation power of GNN, how well does it perform in weather forecast; 2) compare the performance of GNN with traditional ML models; and 3) how do different data preprocessing methods, GNN architectures, and hyperparameters affect the accuracy of weather prediction.

2. Related Works

Numerical weather prediction (NWP) has been the gold standard of weather forecast. It is a branch of scientific computing involving solving complex mathematical equations, such as partial differential equations, with a lot of variables to simulate weather. However, as the method is computationally intensive and induces larger errors when time goes on (Brotzge et al., 2023), some people complement it with machine learning methods.

Deep learning models such as LSTM leverage temporal data to discover trend and give accurate predictions (Li et al., 2023). GNNs (Lam et al., 2022) produce comprehensive representations that captures spatial relationship through message passing mechanism. Generating realistic forecasts with vision transformer (ViT) is another a promising direction (Kurth et al., 2023).

3. Methodology

3.1. Dataset

The experiment leverages a weather dataset retrieved from European Climate Assessment & Dataset (ECA&D) (Klein Tank & Coauthors, 2002; Huber et al., 2024). It contains daily observations at meteorological stations in different parts of Europe. The dataset was constructed of the daily weather observations of 18 stations across 2000-2010. The objective is to train models that predict the mean temperature of the coming day using previous data with spectral and temporal features.

^{*}Equal contribution ¹School of Computing and Data Science, The University of Hong Kong, Hong Kong. Correspondence to: Ambrose Lee <ambr0se@connect.hku.hk>.

3.2. Data Preprocessing

To create the graph, we first calculated the distance between 18 nodes (places), and each region is given an incoming edge with its k-nearest neighbour (default k=3). Then, we generated features for each region. As the set of features differs among regions, we choose features that appear above 80% of the region. For region that does not have the specific feature, we will impute the average value of the feature to the missing column. After filtering less common features, we obtain a unified features consisting of ['sunshine', 'humidity', 'temp_max', 'pressure', 'precipitation', 'temp_min', 'global_radiation', 'temp_mean']. After that, we shuffle the data and split them into train, validation and test set (with a ratio of 6:2:2). For each (x, y) in the dataset, the input x is the graph that contains unified features of 18 regions in the past few days (depends on the temporal window, default temporal_window=3) and edges (with distance as attribute), and the prediction y is the mean temperature of the 18 regions in the coming day.

3.3. Models, Training and Evaluation

In this experiment, we compared the performance of traditional models and GNNs.

For GNNs, we included graph convolutional network (GCN) (Kipf & Welling, 2016), graph attention network (GAT) (Veličković et al., 2017) and GraphSAGE (Hamilton et al., 2017). Here, GCN utilises edge attributes by taking the inverse of each edge attribute (distance) as edge weight. We implemented the GNN architectures and training with PyTorch Geometric.

For traditional ML models, we included linear regression, multi-layer perceptron and random forest. To train the traditional models, we convert the graph features into tabular features, which input is the concatenation of the node attributes and edge attributes of the 18 regions. Additionally, we included a baseline that predicts today's mean temperature with yesterday's mean temperature.

The default setting of the graph is forming incoming edges with the top-3 nearest neighbours, and each node has a temporal windows of 3. For all trainable models, their parameters learnt using Adam optimiser with rate 0.01. The number of layers and batch size of GNNs are 3 and 64 respectively.

To compare the performance of different models, we used mean square error (MSE) as the evaluation metrics. We trained each model for 3 times, and show the mean and standard deviation of the results on the table.

Table 1. Result comparison between models under default setting

| MODELS | TRAIN LOSS | VALIDATION LOSS | TEST LOSS | TIME (IN S) |
|------------|----------------|-----------------|----------------|-------------|
| YESTERDAY | 4.7±0.0 | 4.8±0.0 | 4.7±0.0 | 0.0 |
| GCN | 21.2±0.0 | 21.4±0.2 | 21.2±0.4 | 41.6 |
| GCN_EDGE | 15.9±0.2 | 16.5±0.3 | 16.5±0.4 | 42.2 |
| GAT | 26.9±3.4 | 27.1±7.6 | 27.3±8.0 | 58.5 |
| GAT_EDGE | 24.4±0.0 | 21.2±0.2 | 21.0±0.2 | 59.1 |
| GRAPH_SAGE | 4.6±0.2 | 6.5±1.6 | 6.6±1.7 | 43.9 |
| LR | 4.4±0.0 | 4.7±0.0 | 4.5±0.0 | 0.1 |
| MLP | 5.4±0.1 | 6.0±0.3 | 5.5±0.3 | 1.1 |
| RF | 0.6±0.0 | 4.3±0.0 | 4.1±0.0 | 38.7 |

Table 2. Test loss of models with different k-nearest neighbour

| k | 1 | 3 | 5 | 7 | 10 |
|-----------|-----------------|----------------|----------------|----------|----------------|
| GCN | 16.4±0.4 | 21.1±0.5 | 21.2±0.2 | 22.5±0.1 | 24.5±0.2 |
| GCN_edge | 13.7±0.3 | 14.7±0.3 | 16.2±0.3 | 17.4±0.2 | 18.4±0.3 |
| GAT | 15.9±1.5 | 18.0±4.1 | 18.9±8.1 | 20.3±5.2 | 29.5±4.4 |
| GAT_edge | 13.4±1.8 | 9.0±4.9 | 6.4±3.2 | 8.1±3.3 | 5.3±0.5 |
| GraphSAGE | 6.8±1.4 | 7.8±1.5 | 11.7±1.7 | 8.6±1.2 | 4.7±1.3 |
| LR | 4.4±0.0 | 4.5±0.1 | 4.4±0.0 | 4.5±0.0 | 4.6±0.0 |
| MLP | 5.3±0.4 | 5.2±0.6 | 5.7±0.1 | 5.5±0.3 | 6.1±0.4 |
| RF | 4.1±0.0 | 4.1±0.0 | 4.1±0.0 | 4.1±0.0 | 4.1±0.0 |

4. Experiment & Analysis

4.1. Evaluation on models under default setting

From Table 1, we find that most models performed worse than the baseline model. The baseline model, which predicts current mean temperature with yesterday mean temperature, only has a MSE of 4.7 on testing data. While traditional models like linear regression and random forest got better results (4.5 and 4.1 respectively), most GNNs attained unfavourable results. Except GraphSAGE, other GNNs models had MSE larger than 15 on training, validation and test set, implying they cannot effectively learn the prediction task from the dataset. To improve the result, we experimented with different data preprocessing and training hyperparameters.

4.2. Adjusting preprocessing hyperparameters

We first investigated the impact of k (number of edges/neighbouring nodes per station) on the results. Shown in Table 2, as k increases, most models performed worse except GAT_edge and GraphSAGE. The majority of the models perform the best when there is only 1 neighboring node, while GAT_edge and GraphSAGE achieved the optimum (5.3 and 4.7 respectively) when k=10.

In Table 3, we found that a larger temporal window does not improve the result; conversely, it significantly harmed the accuracy of GAT and GAT_edge.

4.3. Adjusting model and training hyperparameters

Furthermore, we investigated the impacts of batch size and number of layer on accuracy. In Table 4, we identified that a smaller batch size not only increases training time but also

Table 3. Test loss of models with different temporal window

| TEMPORAL WINDOW | 1 | 3 | 5 | 10 | 30 | 50 |
|-----------------|-----------------|-----------------|-----------------|-----------|-----------|-----------|
| GCN | 21.1±0.1 | 20.8±0.2 | 21.3±0.8 | 21.2±0.2 | 21.3±0.3 | 22.2±1.1 |
| GCN_EDGE | 15.3±0.5 | 15.2±0.3 | 15.0±0.5 | 15.7±0.5 | 15.1±0.1 | 15.6±1.2 |
| GAT | 19.5±1.1 | 27.0±5.2 | 35.6±2.8 | 31.6±6.7 | 35.4±3.4 | 48.1±15.6 |
| GAT_EDGE | 7.1±1.4 | 9.8±2.0 | 14.1±4.9 | 30.5±12.7 | 34.9±10.8 | 55.6±19.8 |
| GRAPHSAGE | 5.5±0.8 | 6.0±0.9 | 6.9±1.9 | 8.0±0.9 | 8.6±1.8 | 12.1±1.8 |
| LR | 3.6±0.0 | 3.8±0.0 | 3.9±0.0 | 4.8±0.1 | 5.8±0.1 | 6.8±0.1 |
| MLP | 4.0±0.1 | 4.3±0.2 | 4.2±0.1 | 6.2±1.4 | 7.3±1.7 | 7.6±0.6 |
| RF | 3.8±0.0 | 4.1±0.0 | 4.0±0.0 | 4.2±0.0 | 4.2±0.0 | 4.5±0.0 |

Table 4. Test loss (and seconds needed for training 100 epochs) of GNNs with different batch size

| BATCH SIZE | 1 | 4 | 32 | 64 | 256 |
|------------|-----------|-----------|-----------------|-----------------|----------------|
| GCN | 26.6±0.4 | 23.9±2.6 | 20.9±0.1 | 21.2±0.7 | 21.4±0.6 |
| GCN_EDGE | 25.1±3.2 | 17.8±1.6 | 15.0±0.5 | 14.7±0.2 | 15.3±0.5 |
| GAT | 59.1±17.0 | 59.2±17.0 | 25.7±5.8 | 18.1±4.7 | 20.5±1.3 |
| GAT_EDGE | 57.9±18.7 | 36.4±2.4 | 10.1±1.0 | 6.5±1.5 | 9.7±3.8 |
| GRAPHSAGE | 10.3±2.4 | 7.0±0.6 | 6.7±1.8 | 7.2±0.3 | 6.3±0.7 |

Table 5. Test loss of GNNs with different number of layers

| # OF LAYERS | 1 | 2 | 3 | 4 | 5 |
|-------------|-----------------|----------------|----------|----------|----------|
| GCN | 18.1±0.2 | 19.9±0.4 | 20.7±0.1 | 21.2±0.2 | 21.7±0.2 |
| GCN_EDGE | 9.4±0.1 | 13.2±0.2 | 14.7±0.1 | 16.1±0.6 | 17.6±0.7 |
| GAT | 12.6±4.1 | 19.7±1.5 | 19.0±3.4 | 22.9±1.8 | 22.4±0.4 |
| GAT_EDGE | 11.5±2.3 | 5.3±0.3 | 7.1±3.3 | 11.4±7.1 | 45.6±6.0 |
| GRAPHSAGE | 4.0±0.0 | 3.9±0.3 | 7.7±1.9 | 13.6±0.4 | 6.6±1.2 |

harm the results. Most models reached the best and stable results when batch size is larger than or equal to 32.

Also, more layers does not necessarily lead to better results. In this simple task, most GNNs achieved the optimal results when the model is only consist of 1-2 layer(s) according to Table 5.

Note, considering edge attributes (like GCN_edge and GAT_edge) could give models extra information and improve the results. More detailed results can be found in the ‘TemperatureForecast/result_comparison’.

4.4. Theoretical Analysis

As GNN can take in any graph structure, the existing model can still be used even when new edges or nodes are added to the graph; on the other hand, the graph structure and the sequence of the input attributes needed to be the same in order to reuse traditional ML models.

However, GNNs have a higher computational cost compared to traditional models, in turn taking longer time and more compute to train (as shown in Table 1).

5. Limitation & Future Work

Although we have included 8 features for each region at each timestamp, the unsatisfactory results of GNNs indicates that the trained models failed to generate a good representation

of the weather of each region. Adding more attributes with domain knowledge, including geospatial feature, occurrence of extreme events, cloud cover and time could largely enhance the representation of the model. Normalisation could also help the model to learn a better representation.

Besides, as historical data and past trends were proven to have significant impact on future temperature, using more advanced feature engineering techniques (such as moving average of different features across time) or kernel (such as LSTM, RNN or Transformer) to handle time-series data could better capture long-term dependencies and yield a better prediction.

Moreover, although the geographical location of the stations possess graph structure on the map, the distances between stations are large. Thus, edges might not signifies strong relationships and neighbouring nodes might not have significant impact on the prediction, hindering the performance of GNNs.

In addition, observing the graphs that visualise loss over epochs (can be found in TemperatureForecast/result_comparison), most GNNs started to converge before 30th epochs. Instead of full training with every hyperparameter set (100 iterations), it is better to do cross validation to find the best set of hyperparameters with small number of iterations (like 20-30), and do full training after finding the best set of hyperparameters.

6. Conclusion

In this paper, we have examined the accuracy of predicting tomorrow’s temperature of different models, and compared the impacts of different hyperparameters. The results indicated that not all GNNs perform well on graph-structured data, especially when edges do not signify important and relevant relations. Besides, the performance of the models were limited by the range and relevancy of features. Lastly, this experiment has reminded us that despite the emergence of advanced architectures, data preprocessing, feature engineering and model selection remain essential to train accurate machine learning models.

References

- Brotzge, J. A., Berchhoff, D., Carlis, D. L., Carr, F. H., Carr, R. H., Gerth, J. J., Gross, B. D., Hamill, T. M., Haupt, S. E., Jacobs, N., McGovern, A., Stensrud, D. J., Sztankowski, G., Szunyogh, I., and Wang, X. Challenges and opportunities in Numerical weather prediction. *Bulletin of the American Meteorological Society*, 104(3):E698–E705, 2 2023. doi: 10.1175/bams-d-22-0172.1. URL <https://journals.ametsoc.org/view/journals/bams/104/3/BAMS-D-22-0172.1.xml>.
- Hamilton, W. L., Ying, R., and Leskovec, J. Inductive representation learning on large graphs, 6 2017. URL <https://arxiv.org/abs/1706.02216>.
- Huber, F., van Kuppevelt, D., Steinbach, P., Sauze, C., Liu, Y., and Weel, B. Will the sun shine? – an accessible dataset for teaching machine learning and deep learning, 2024. URL <https://proceedings.mlr.press/v207/huber23a/huber23a.pdf>.
- Kipf, T. N. and Welling, M. Semi-Supervised Classification with Graph Convolutional Networks, 9 2016. URL <https://arxiv.org/abs/1609.02907>.
- Klein Tank, A. and Coauthors. Daily dataset of 20th-century surface air temperature and precipitation series for the european climate assessment. *International Journal of Climatology*, 22:1441–1453, 2002. URL <http://www.ecad.eu>.
- Kurth, T., Subramanian, S., Harrington, P., Pathak, J., Mardani, M., Hall, D., Miele, A., Kashinath, K., and Anandkumar, A. FourCastNet: Accelerating Global High-Resolution Weather Forecasting Using Adaptive Fourier Neural Operators. *association for computing machinery*, pp. 1–11, 6 2023. doi: 10.1145/3592979.3593412. URL <https://doi.org/10.1145/3592979.3593412>.
- Lam, R., Sanchez-Gonzalez, A., Willson, M., Wirnsberger, P., Fortunato, M., Alet, F., Ravuri, S., Ewalds, T., Eaton-Rosen, Z., Hu, W., Merose, A., Hoyer, S., Holland, G., Vinyals, O., Stott, J., Pritzel, A., Mohamed, S., and Battaglia, P. GraphCast: Learning skillful medium-range global weather forecasting, 12 2022. URL <https://arxiv.org/abs/2212.12794>.
- Li, Y., Li, T., Lv, W., Liang, Z., and Wang, J. Prediction of daily temperature based on the robust machine learning algorithms. *Sustainability*, 15(12):9289, 6 2023. doi: 10.3390/su15129289. URL <https://www.mdpi.com/2071-1050/15/12/9289>.
- Lobell, D. B. and Field, C. B. Global scale climate–crop yield relationships and the impacts of recent warming. *Environmental Research Letters*, 2(1):014002, 3 2007. doi: 10.1088/1748-9326/2/1/014002. URL <https://doi.org/10.1088/1748-9326/2/1/014002>.
- Veličković, P., Cucurull, G., Casanova, A., Romero, A., Liò, P., and Bengio, Y. Graph attention networks, 10 2017. URL <https://arxiv.org/abs/1710.10903>.
- Xu, K., Hu, W., Leskovec, J., and Jegelka, S. How Powerful are Graph Neural Networks?, 10 2018. URL <https://arxiv.org/abs/1810.00826>.
- Zhou, J., Cui, G., Hu, S., Zhang, Z., Yang, C., Liu, Z., Wang, L., Li, C., and Sun, M. Graph Neural Networks: A Review of Methods and applications, 12 2018. URL <https://arxiv.org/abs/1812.08434>.